

## CS186: Introduction to Databases

1. Lecture 01 - Introduction
2. Lecture 02 - Sorting and Hashing
3. Lecture 03 - Single-Table Queries
4. Lecture 04 - Join Algorithms
5. Lecture 05 - Sorting Data: Disks and Files
6. Lecture 06 - File Organizations and Indexing
7. Lecture 07 - Tree-Structured Indexes
8. Lecture 08 - Relational Algebra
9. Lecture 09 - SQL: The Query Language
10. Lecture 10 - SQL: Constraints, SQL Embedded in Other Languages
11. Lecture 11 - SQL (Cont.), Entity-Relationship Diagrams and the Relational Model
12. Lecture 12 - Entity-Relationship Diagrams and the Relational Model (Cont.)
13. Lecture 13 - Functional Dependencies and Schema Refinement
14. Lecture 14 - Functional Dependencies and Schema Refinement
15. Lecture 15 - Relational Query Optimization I
16. Lecture 16 - Relational Query Optimization II
17. Lecture 17 - Elementary IR: Scalable Boolean Text Search
18. Lecture 18 - Ranking and Crawling, Data Visualization
19. Lecture 19 - Data Visualization (Cont.), Transactions and Concurrency Control
20. Lecture 20 - Transactions and Concurrency Control (Cont.)
21. Lecture 21 - Big Data Analytics Systems (Guest Lecture by Reynold Xin)
22. Lecture 22 - Advanced Topics in Concurrency Control
23. Lecture 23 - Advanced Topics in Concurrency Control (Cont.), Crash Recovery
24. Lecture 24 - A Perspective on Data Science (Guest Lecture by Tye Rattenbury)
25. Lecture 25 - Crash Recovery (Cont.)
26. Lecture 26 - Crash Recovery Wrap-up; Distributed Data, Replication and NoSQL

- 27. Lecture 27 - Review I
  - 28. Lecture 28 - Review II
- 

- 1. 1
  - 1.1. Lecture 1 Part 1 Intro and Why
  - 1.2. Lecture 1 Part 2 What
  - 1.3. Lecture 1 Part 3 Who
  - 1.4. Lecture 1 Part 4 How
- 2. 2
- 3. 3
- 4. 4
  - 4.1. Lec 4 Part 1 Architecture
  - 4.2. Lec 4 Part 2 Storage Hierarchy
  - 4.3. Lec 4 Part 3 Components of a Disk
  - 4.4. Lec 4 Part 4 Flash
  - 4.5. Lec 4 Part 5 Storage Pragmatics and Trends
  - 4.6. Lec 4 Part 6 Disk Space Management
- 5. 4.5
  - 5.1. Lec 4 5 Part 1 Files
  - 5.2. Lec 4 5 Part 2 Pages for Fixed Length
  - 5.3. Lec 4 5 Part 3 Pages for Variable Length
  - 5.4. Lec 4 5 Part 4 Record Layout
  - 5.5. Lec 4 5 Part 5 Summary
- 6. 5
  - 6.1. Lecture 5 Part 1 Review and Goals
  - 6.2. Lec 5 Part 2 Model and Assumptions
  - 6.3. Lec 5 Part 3 Scan
  - 6.4. Lec 5 Part 4 Equality
  - 6.5. Lec 5 Part 5 Range Search
  - 6.6. Lec 5 Part 6 Insert
  - 6.7. Lec 5 Part 7 Delete
  - 6.8. Lec 5 Part 8 Summary
- 7. 6
  - 7.1. Lec 6 Part 1 Intro to Indexes
  - 7.2. Lec 6 Part 2 High Fan out Search Tree
  - 7.3. Lec 6 Part 3 Search and Insertion in ISAM
  - 7.4. Lecture 6 Part 4 B+ Tree Intro
  - 7.5. Lecture 6 Part 5

- 7.6. Lecture 6 Part 6
- 7.7. Lecture 6 Part 7
- 7.8. Lecture 6 Part 8 B+ Tree Deletion
- 7.9. Lecture 6 Part 9
- 7.10. Lecture 6 Part 10
- 8. 7
  - 8.1. Lec 7 Part 1 Lecture Introduction
  - 8.2. Lec 7 Part 2 Query Support Overview
  - 8.3. Lec 7 Part 3 Composite Search Keys in One Dimension
  - 8.4. Lec 7 Part 4 Three Alternatives for Data Entry Storage
  - 8.5. Lec 7 Part 5 Clustered vs Unclustered Indexes
  - 8.6. Lec 7 Part 6 Variable Length Keys: Occupancy and Compression
  - 8.7. Lec 7 Part 7 Clustered B+ Tree Cost Model, Lec 7 Part 7 Clustered B+ Tree Cost Model \*\*REVISED\*\*
  - 8.8. Lec 7 Part 8 Lecture Summary
- 9. 8
  - 9.1. Lec 8 Part 1 Intro to Buffer Management
  - 9.2. Lec 8 Part 2 Mapping Pages
  - 9.3. Lec 8 Part 3 Dirty Pages
  - 9.4. Lec 8 Part 4 Buf Mgr State
  - 9.5. Lec 8 Part 5 Page Replacement
  - 9.6. Lec 8 Part 6 LRU
  - 9.7. Lec 8 Part 7 Clock
  - 9.8. Lec 8 Part 8 LRU and Sequential Flooding
  - 9.9. Lec 8 Part 9 MRU and Repeated Scans
  - 9.10. Lec8 Part 10 Additional Topics
  - 9.11. Lec 8 Part 11 Summary
- 10. 9
  - 10.1. Lec 9 Part 1 Out of Core Algs
  - 10.2. Lec 8 Part 2 Mapping Pages
  - 10.3. Lec 9 Part 3 Sort Hash Specs
  - 10.4. Lec 9 Part 4 Two Way Sorting
  - 10.5. Lecture 9 Part 5 External Sort
  - 10.6. Lec 9 Part 6 External Hashing
  - 10.7. Lec 9 Part 7 Recursive Partitioning for External Hash
  - 10.8. Lec 9 Part 8 Sort Hash Duality
  - 10.9. Lec 9 Part 9 parallel hashing
  - 10.10. Lec 9 Part 10 parallel sort
  - 10.11. Lec 9 Part 11 Sort v Hash
  - 10.12. Lec 9 Part 12 Summary
- 11. 10
  - 11.1. Lec 10 Part 1 Intro

- 11.2. Lec 10 Part 2 Unary Operators
- 11.3. Lec 10 Part 3 Set Operators
- 11.4. Lec 10 Part 4 Joins
- 11.5. Lec 10 Part 5 GroupBy and Summary
- 12. 12
  - 12.1. Lecture 11 Part 1 Intro and Select
  - 12.2. Lecture 11 Part 2 HeapScan and Sort Iterators
  - 12.3. Lecture 11 Part 3 Group By Iterator
  - 12.4. Lecture 11 Part 4 Cost Model for Subsequent Analysis
  - 12.5. Lecture 11 Part 5 Nested Loops Join
  - 12.6. Lecture 11 Part 6 Index Nested Loops Join
  - 12.7. Lecture 11 Part 7 Sort Merge Join
  - 12.8. Lecture 11 Part 8 Sort Merge Refinement, Lecture 11 Part 8 Sort Merge Cost and a Refinement
  - 12.9. Lec 11 Part 9
  - 12.10. Lecture 11 Part 10, Lecture 11 Part 10 Grace Hash Animation
  - 12.11. Lecture 11 Part 11 Hash Join Costs
  - 12.12. Lecture 11 Part 12 Conclusion
- 13. 13
  - 13.1. Lecture 13 Part 1 Intro
  - 13.2. Lecture 13 Part 2 Big Picture
  - 13.3. Lecture 13 Part 3 Algebra Equivalences
  - 13.4. Lecture 13 Part 4 Heuristics
  - 13.5. Lecture 13 Part 5 Physical Equivalences
  - 13.6. Lecture 13 Part 6 Running Example
  - 13.7. Lecture 13 Part 7 Selection Pushdown
  - 13.8. Lecture 13 Part 8 Join Orders
  - 13.9. Lecture 13 Part 9 Join Algs
  - 13.10. Lecture 13 Part 10 Projection Pushdown
  - 13.11. Lecture 13 Part 11 Indexes
  - 13.12. Lecture 13 Part 12 Conclusion
- 14. 14
  - 14.1. Lecture 14 Part 1 Introduction
  - 14.2. Lecture 14 Part 2 Background Query Blocks and Physical Properties
  - 14.3. Lecture 14 Part 3 Plan Space
  - 14.4. Lecture 14 Part 4 Selectivity Estimation
  - 14.5. Lecture 14 Part 5 Selectivity in More Depth
  - 14.6. Lecture 14 Part 6 Using Histograms for Selectivity Estimation
  - 14.7. Lecture 14 Part 7 Selectivity for Join Queries
  - 14.8. Lecture 14 Part 8 Summary of Cost Estimation
  - 14.9. Lecture 14 Part 9 Single Table Plans
  - 14.10. Lecture 14 Part 10 Dynamic Programming
  - 14.11. Lecture 14 Part 11 Example Query

- 14.12. Lecture 14 Part 12 Summary of Query Optimization
- 15. 15
  - 15.1. Lecture 15 Part 1 Intro to Text Search
  - 15.2. Lecture 15 Part 2 Bag of Words
  - 15.3. Lecture 15 Part 3 Inverted Files and Single-Term Queries
  - 15.4. Lecture 15 Part 4 Inverted Indexes and Boolean Logic
  - 15.5. Lecture 15 Part 5 Phrase Queries and Content Results
  - 15.6. Lecture 15 Part 6 Updating Text Indexes, Lecture 15 Part 6 Updating Text Indexes
  - 15.7. Lecture 15 Part 7 Text Search Summary
- 16. 16
  - 16.1. Lecture 16 Introduction
  - 16.2. Lecture 16 Part 2 Overview of DB Design
  - 16.3. Lecture 16 Part 3 Data Models and Relational Levels of Abstraction
  - 16.4. Lecture 16 Part 4 Data Independence
  - 16.5. Main Sub 05
  - 16.6. Lecture 16 Part 6 Basics of Entities and Relationships
  - 16.7. Lecture 16 Part 7 Key and Participation Constraints
  - 16.8. Lecture 16 Part 8 Weak Entities
  - 16.9. Lecture 16 Part 9 Alternative Notation and Terminology
  - 16.10. Lecture 16 Part 10 Binary vs Ternary Relationships
  - 16.11. Lecture 16 Part 11 Aggregation and Ternary Relationships
  - 16.12. Lecture 16 Part 12 Entities vs Attributes
  - 16.13. Lecture 16 Part 13 Entities vs Relationships
  - 16.14. Lecture 16 Part 14 Converting ER to Relational
  - 16.15. Lecture 16 Part 15 Translating Weak Entity Sets
  - 16.16. Lecture 16 part 16 Summary
- 17. 17
  - 17.1. Lecture 17 Part 1 Functional Dependencies
  - 17.2. Lecture 17 Part 2 Anomalies
  - 17.3. Lecture 17 Part 3 Armstrongs Axioms
  - 17.4. Lecture 17 Part 4 Attribute Closures
  - 17.5. Lecture 17 Part 5 BCNF and other Normal Forms
  - 17.6. Lecture 17 Part 6 Lossless Join Decompositions
  - 17.7. Lecture 17 Part 7 Dependency Preservation and BCNF Decomposition
  - 17.8. Lecture 17 Part 8 Conclusion
- 18. 18
  - 18.1. Lec 18 Part 1 Intro to Transactions
  - 18.2. Lec 18 Part 2 Examples and Problems
  - 18.3. Lec 18 Part 3 The Transaction Model
  - 18.4. Lect 18 Part 4 ACID Properties of Transactions
  - 18.5. Lec 18 Part 5 Serializability
  - 18.6. Lec 18 Part 6 Conflict Serializability

- 18.7. Lec 18 Part 7 Conflict Dependency Graphs
- 18.8. Lecture 18 Part 8 View Serializability
- 19. 19
  - 19.1. Lecture 19 Part 01 Two Phase Locking
  - 19.2. Lecture 19 Part 2 Strict 2PL
  - 19.3. Lecture 19 Part 3 Concurrency Examples
  - 19.4. Lecture 19 Part 04 Lock Manager Implementation
  - 19.5. Lecture 19 Part 5 Deadlock Scenarios and Prevention
  - 19.6. Lecture 19 Part 6 Deadlock Avoidance
  - 19.7. Lecture 19 Part 7 Deadlock Detection and Resolution
  - 19.8. Lecture 19 Part 8 Intro to Multigranularity LockingLecture 19 Part 9 Intent Locks and Hierarchical Lock Protocol
  - 19.9. Lecture 19 Part 9 Intent Locks and Hierarchical Lock Protocol
  - 19.10. Lecture 19 Part 10 FYI: Index Latching and Phantoms
  - 19.11. Lecture 19 Part 11 Conclusion
- 20. 20
  - 20.1. Lecture 20 Part 1 Recovery Intro
  - 20.2. Lecture 20 Part 2 Need for Atomicity and Durability, SQL support for Transactions
  - 20.3. Lecture 20 Part 3 Strawman Solution
  - 20.4. Lecture 20 Part 4 STEAL / NO FORCE, UNDO and REDO
  - 20.5. Lecture 20 Part 5 Intro to Write-Ahead Logging
  - 20.6. Lecture 20 Part 6: ARIES Logging
  - 20.7. Lecture 20 Part 7: ARIES Logging Simple Case
  - 20.8. Lecture 20 Part 8 ARIES Abort and Checkpointing
  - 20.9. Lecture 20 Part 9 ARIES Recovery: Overview and Analysis Phase
  - 20.10. Lecture 20 Part 10: ARIES Recovery, REDO Phase
  - 20.11. Lecture 20 Part 11: ARIES Recovery, UNDO Phase
  - 20.12. Lecture 20 Part 12: Example Scenarios, Additional FAQs
  - 20.13. Lecture 20 Part 13: Recovery Conclusion
- 21. 21
  - 21.1. Lecture 21 1 Cosine Similarity
  - 21.2. Lecture 21 2 TFxIDF
  - 21.3. Lecture 21 3 Precision and Recall
  - 21.4. Lecture 21 4 Parallelization and Zipf Distributions
  - 21.5. Lecture 21 5 n-grams, q-grams and other tricks
  - 21.6. Lecture 21 6 PageRank
  - 21.7. Lecture 21 7 Notes from the Real World
  - 21.8. Lecture 21 8 Webcrawler Intro
  - 21.9. Lecture 21 9 Webcrawler Refinements
  - 21.10. Lecture 21 10 Conclusion
- 22. 22
  - 22.1. Lecture 22 Part 1 Introduction to Distributed Transactions

- 22.2. Lecture 22 Part 2 Distributed Locking
- 22.3. Lecture 22 Part 3 Distributed Deadlock Detection
- 22.4. Lecture 22 Part 4 Intro to Distributed Commit
- 22.5. Lecture 22 Part 5 Simple 2PC: Messaging Only
- 22.6. Lecture 22 Part 6 Full 2PC Including Logging (Presumed Abort)
- 22.7. Lecture 22 Part 7 The Recovery Processes
- 22.8. Lecture 22 Part 8 2PC, Locking and Availability
- 22.9. Lecture 22 Part 9: Summary of Distributed Transactions
- 23. SQL1
  - 23.1. SQL 1 Clip 1
  - 23.2. SQL 1 Clip 2
  - 23.3. SQL 1 Clip 3
  - 23.4. SQL 1 Clip 4
  - 23.5. SQL 1 Clip 5
  - 23.6. SQL 1 Clip 6 7
- 24. SQL2
  - 24.1. SQL II 1 4
  - 24.2. SQL II 6 10
  - 24.3. SQL II 11 12
  - 24.4. SQL II 13 15 slides
  - 24.5. SQL II 17 21
  - 24.6. SQL II 23 28
  - 24.7. SQL II 30 34
  - 24.8. SQL II 35
  - 24.9. SQL II 37 39
  - 24.10. SQL II 41 46
  - 24.11. SQL II 48 53
  - 24.12. SQL II 55 56
  - 24.13. SQL II 58 66
  - 24.14. SQL II 68 69
- 25. Misc
  - 25.1. A Bit of History
  - 25.2. Kinds of Query Parallelism
  - 25.3. Parallel Hash Join
  - 25.4. Parallel Grouping/Aggregation
  - 25.5. One-Sided and Broadcast Join Patterns
  - 25.6. Symmetric Hash Join
  - 25.7. Intro to Parallelism
  - 25.8. Parallel Architectures and Software Structures
  - 25.9. Parallel Data Access
  - 25.10. Parallel Sort-Merge Join
  - 25.11. Summary of Parallel Query Execution